

SQL Survival Guide

Joins

Andy Choens

EBCoP

May 18, 2015

Outline

- 1 Example Data
- 2 Dafynitions
- 3 SQL Order Of Operations
- 4 ANSI Joins
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins
- 6 Questions?

Outline for section 1

- 1 Example Data
- 2 Dafynitions
- 3 SQL Order Of Operations
- 4 ANSI Joins
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins
- 6 Questions?

Example Data

- Two tables of example data!
- ALL Code provided: `sql/04-joins`
 - ▶ `create_tables.sql`: Creates the example tables.
 - ▶ Everything Else: Other files contain example queries.
- Example data in SQLite: `data/04-joins.sqlite`.

Table: Departments

DEPT_ID	DEPT_NAME	DEPT_FLOOR
31	Sales	1
33	Engineering	3
34	Clerical	2
35	Marketing	3

Aliased as 'dept'.

Table: Employees

EID	DEPT_ID	LAST_NAME	FIRST_NAME	GENDER
1	31	Rafferty	Gerry	M
3	33	Jones	Jon	M
5	33	Heisenberg	Werner	M
7	34	Robinson	Elizabeth	F
9	34	Smith	Jefferson	M
11	NULL	Williams	Serena	F

Aliased as 'empl'.

Outline for section 2

- 1 Example Data
- 2 Dafynitions**
- 3 SQL Order Of Operations
- 4 ANSI Joins
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins
- 6 Questions?

Dafynitions?

We need a common set of terms to discuss this.

TODO: The slides in this section need a lot of work.

Dafynition: Relational Data

TODO

Dafynition: Normalized Data

TODO

Dafynition: Result Set

TODO

Dafynition: Clause

The following protected SQL commands START a clause:

- SELECT
- FROM
- WHERE
- HAVING

Dafynition: Clause

The following protected SQL commands START a clause:

- SELECT
- FROM
- WHERE
- HAVING

Style:

- Align SQL Clauses
- Predicates should be indented consistently.
- Feel free to ask me why I structure SQL queries the way I do.

Dafynition: Join Clause?

- The FIRST thing you should write.
- Combines records from two or more tables (result set).
- SQL Joins are a difficult skill to master.
- They are necessary for working with normalized, relational data.

Outline for section 3

- 1 Example Data
- 2 Dafynitions
- 3 SQL Order Of Operations**
- 4 ANSI Joins
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins
- 6 Questions?

Order of Operations

What is the answer to this (silly) equation?

$$\sqrt{(2^2 + 2) \cdot 6}$$

Order of Operations

What is the answer to this (silly) equation?

$$\sqrt{(2^2 + 2) \cdot 6}$$

And how were we all able to come up with the same answer?

Order of Operations

Just like math, SQL has an order of operations:

- 1 FROM
- 2 WHERE
- 3 GROUP BY
- 4 SELECT
- 5 HAVING
- 6 ORDER BY

Order of Operations

Some last order of operation notes:

- Sub-queries are run before the outer query.
- The optimizer may reorganize query (relational algebra).
- Sometimes writing things out of order can bite you.

Outline for section 4

- 1 Example Data
- 2 Dafynitions
- 3 SQL Order Of Operations
- 4 ANSI Joins**
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins
- 6 Questions?

Types of Joins

There are **only** 5 ANSI joins:

- | | |
|--------------------|------------------------------------|
| ① Cross Join | ① DO NOT USE! |
| ② Inner Join | ② Very Useful. |
| ③ Left Outer Join | ③ Very Useful. |
| ④ Right Outer Join | ④ Most of you should not use this. |
| ⑤ Full Outer Join | ⑤ You will rarely use this. |

Most SQL queries can be written with only 2 types of joins.

Cross Join: Discussion

- Returns the Cartesian product of the tables in the FROM clause.
- Can be written explicitly and implicitly.
- This is **ALMOST ALWAYS** a bad idea.

Explicit Cross Join: Example

Question:

How many rows will the following query return?

```
1  select *
2  from EMPLOYEES empl cross join DEPARTMENTS dept
3  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/cross-joins.sql>

Explicit Cross Join: Example

Question:

How many rows will the following query return?

```
1  select *
2  from EMPLOYEES emp1 cross join DEPARTMENTS dept
3  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/cross-joins.sql>

Answer:

24 rows

Cross Join: What Happens

TODO: A graphical representation of what it is doing.

Cross Join: Result Set

EID	DEPT_ID	LAST_NAME	FIRST_NAME	GENDER	DEPT_ID	DEPT_NAME	DEPT_FLOOR
1	31	Rafferty	Gerry	M	31	Sales	1
1	31	Rafferty	Gerry	M	33	Engineering	3
1	31	Rafferty	Gerry	M	34	Clerical	2
1	31	Rafferty	Gerry	M	35	Marketing	3
3	33	Jones	Jon	M	31	Sales	1
3	33	Jones	Jon	M	33	Engineering	3
3	33	Jones	Jon	M	34	Clerical	2
3	33	Jones	Jon	M	35	Marketing	3
5	33	Heisenberg	Werner	M	31	Sales	1
5	33	Heisenberg	Werner	M	33	Engineering	3
5	33	Heisenberg	Werner	M	34	Clerical	2
5	33	Heisenberg	Werner	M	35	Marketing	3
7	34	Robinson	Elizabeth	F	31	Sales	1
7	34	Robinson	Elizabeth	F	33	Engineering	3
7	34	Robinson	Elizabeth	F	34	Clerical	2
7	34	Robinson	Elizabeth	F	35	Marketing	3
9	34	Smith	Jefferson	M	31	Sales	1
9	34	Smith	Jefferson	M	33	Engineering	3
9	34	Smith	Jefferson	M	34	Clerical	2
9	34	Smith	Jefferson	M	35	Marketing	3
11	[NULL]	Williams	Serena	F	31	Sales	1
11	[NULL]	Williams	Serena	F	33	Engineering	3
11	[NULL]	Williams	Serena	F	34	Clerical	2
11	[NULL]	Williams	Serena	F	35	Marketing	3

Implicit Cross Join: Example

```
1  select *  
2  from EMPLOYEES empl, DEPARTMENTS dept  
3  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/cross-joins.sql>

- Also returns 24 rows.
- What does this query look like?

Cross Join: Results

- Cross Joins Are Dangerous! (Especially the implicit ones.)
- They return the maximum number of rows possible.

Inner Join: Discussion

- Returns all records which have matching records in both tables, according to the join-predicate or WHERE clause.
- Can be written explicitly or implicitly.
- Returns the least number of rows.

Explicit Inner Join: Example

Question:

How many rows will the following query return?

```
1  select *
2  from EMPLOYEES empl inner join DEPARTMENTS dept
3  on empl.dept_id = dept.dept_id
4  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/cross-joins.sql>

Explicit Inner Join: Example

Question:

How many rows will the following query return?

```
1  select *
2  from EMPLOYEES empl inner join DEPARTMENTS dept
3  on empl.dept_id = dept.dept_id
4  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/cross-joins.sql>

Answer:

5 rows

Inner Join: What Happens

TODO: A pictorial representation.

Inner Join: Result Set

EID	DEPT_ID	LAST_NAME	FIRST_NAME	GENDER	DEPT_ID	DEPT_NAME	DEPT_FLOOR
1	31	Rafferty	Gerry	M	31	Sales	1
3	33	Jones	Jon	M	33	Engineering	3
5	33	Heisenberg	Werner	M	33	Engineering	3
7	34	Robinson	Elizabeth	F	34	Clerical	2
9	34	Smith	Jefferson	M	34	Clerical	2

Implicit Inner Join: Example

This should look familiar:

```
1  select *
2  from EMPLOYEES empl, DEPARTMENTS dept
3  where empl.dept_id = dept.dept_id
4  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/cross-joins.sql>

- Also returns 5 rows.

Why Not Use Implicit Join Syntax?

Question:

Implicit Cross Join v Implicit INNER JOIN: What's the difference?

Why Not Use Implicit Join Syntax?

Question:

Implicit Cross Join v Implicit INNER JOIN: What's the difference?

Answer:

The WHERE clause.

- Implicit Join syntax *is* deprecated.
- This style makes it too easy to write a Cross Join (cartesian).
- It makes it harder to learn the SQL Order of Operations.

Outer Joins

- The result set from an Inner Joins includes matching records ONLY.
- The result set from an Outer Join retains more records.
- Types of Outer Joins:
 - ▶ Left Outer Join (Left Join)
 - ▶ Right Outer Join (Right Join)
 - ▶ Full Outer Join
- The only difference is which records become part of the result set.

Left Outer Join: Discussion

- Result set includes all members of the 'left' table.

Left Outer Join: Example

Question:

How many rows will the following query return?

```
1  select *
2  from EMPLOYEES empl left join DEPARTMENTS dept
3  on empl.department_id = dept_dept_id
4  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/left-joins.sql>

Left Outer Join: Example

Question:

How many rows will the following query return?

```
1  select *
2  from EMPLOYEES empl left join DEPARTMENTS dept
3  on empl.department_id = dept_dept_id
4  ;
```

Source: <https://github.com/Choens/sql-survival-guide/blob/master/sql/04-joins/left-joins.sql>

Answer:

6 rows

Right Outer Join: Discussion

TODO

Right Outer Join: Example

TODO

Full Outer Join: Discussion

TODO

Full Outer Join: Example

TODO

Outline for section 5

- 1 Example Data
- 2 Dafynitions
- 3 SQL Order Of Operations
- 4 ANSI Joins
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins**
- 6 Questions?

Self Join: Discussion

TODO

Self Join: Example

TODO

Outline for section 6

- 1 Example Data
- 2 Dafynitions
- 3 SQL Order Of Operations
- 4 ANSI Joins
 - Cross Joins
 - Inner Joins
 - Left Outer Joins
 - Right Outer Joins
 - Full Outer Joins
- 5 Self Joins
- 6 Questions?

Additional Information

- [https://en.wikipedia.org/wiki/Join_\(SQL\)](https://en.wikipedia.org/wiki/Join_(SQL))

Questions?